

CAROL WILLIAMS

Senior Data Scientist

carol.williams@email.com | New York, NY | +1-555-0103

SUMMARY

Data Scientist with 7 years of experience in machine learning, statistical analysis, and data engineering. Expertise in building production ML models for healthcare and fintech applications. Published 3 papers in top-tier ML conferences.

EXPERIENCE

Senior Data Scientist · HealthAI Corp (2021 - Present)

- Led development of an NLP system for medical record analysis using BERT and GPT
- Built a patient risk prediction model achieving 92% AUC-ROC score
- Designed A/B testing frameworks for evaluating model performance in production
- Managed a team of 4 data scientists and 2 ML engineers
- Reduced model inference latency by 60% through ONNX optimization and model distillation

Data Scientist · FinTech Analytics (2018 - 2021)

- Developed credit scoring models using XGBoost and LightGBM
- Built real-time fraud detection system processing 1M+ transactions/day
- Created customer segmentation models using K-means and DBSCAN clustering
- Automated feature engineering pipeline reducing model development time by 50%
- Presented ML insights to C-suite stakeholders and board members

ML Research Intern · Stanford AI Lab (2017 - 2018)

- Conducted research on transformer architectures for time-series forecasting
- Implemented novel attention mechanisms improving forecast accuracy by 15%
- Co-authored paper accepted at NeurIPS 2018

TECHNICAL SKILLS

- Languages: Python, R, SQL, Julia
- ML/DL: scikit-learn, TensorFlow, PyTorch, Hugging Face, XGBoost, LightGBM
- NLP: spaCy, NLTK, transformers, LangChain
- Data: pandas, NumPy, SciPy, Polars, Apache Spark
- Visualization: Matplotlib, Seaborn, Plotly, Tableau
- MLOps: MLflow, Docker, SageMaker, Kubeflow, DVC
- Tools: Git, Jupyter, Airflow, dbt

EDUCATION

Ph.D. Statistics · Stanford University (2018)

M.S. Data Science · Columbia University (2016)

B.S. Mathematics · MIT (2014)